

RELTWIN: CONTRASTING LOCALLY VALID TIMESTAMP ANSWERS FOR QUERY-SPECIFIC AUDIO GROUNDING

Yuehan Zhang Zhicheng Tang Wei Xu*

Huazhong University of Science and Technology, Wuhan, China
yuehanzhang2005@hotmail.com, u202414252@hust.edu.cn

ABSTRACT

Audio temporal grounding must distinguish which event occurrence answers a query, even when opposite event orders coexist in one recording. Models can instead return the same window for inverse queries, revealing a query-window binding error. We introduce RelTwin, an add-on fine-tuning method that constructs opposite-order windows from shared audio excerpts and contrasts their timestamp answers under each query. Both windows have acoustic support; the requested order determines the correct answer. This answer-level competition complements sequence supervision and preserves direct timestamp generation at inference. Applied to SpotSound-A, RelTwin reaches 59.43, 86.68, and 74.03 mIoU on SpotSound-Bench, Clotho-Moment, and UnAV-100, respectively. Against matched paired SFT, it improves joint localization by 31.25 points on synthetic relation-development data and SpotSound mIoU by 1.02 points. The binding advantage persists under timing intervention, supporting explicit supervision of which occurrence answers the query.

Index Terms— Audio temporal grounding, timestamp answer contrasts, query-window binding, paired supervision

1. INTRODUCTION

Audio temporal grounding maps a query to event intervals [1, 2]. In a recording containing both “dog then rooster” and “rooster then dog,” the descriptions refer to different windows. Recognizing the categories establishes presence; answering the query requires selecting the requested order. Returning the same real window for both queries is a *query-window binding error*. We target this distinction between event detection and query-specific localization.

RelTwin (Fig. 1) addresses this binding problem through *locally valid timestamp contrasts*. Shared excerpts form opposite-order windows in one recording. Each window contains a real event sequence, but its correctness as an answer depends on the query. Sequence supervision rewards correct timestamp answers individually; the added candidate loss

explicitly compares the two answers under each query. Reversing the query exchanges their labels while keeping the recording fixed, directly training the model to choose the requested occurrence.

RelTwin fine-tunes an existing grounding model using its own timestamp-answer likelihoods, preserving direct generation at inference. We evaluate the adapted SpotSound-A model on three public benchmarks and examine answer-level competition through matched training comparisons.

We make three contributions. First, we characterize query-window binding errors when inverse relations co-occur in one recording. Second, we introduce locally valid timestamp contrasts as an additional fine-tuning objective for query-specific answer selection. Third, public grounding evaluations, paired diagnosis, and a timing intervention connect the observed localization performance to improved query-window binding.

2. RELATED WORK

Audio temporal grounding. Weakly supervised grounding learns temporal correspondences from clip-level captions, reducing dependence on frame-level annotations [3]. Audio-language models such as Pengi [4], Qwen2-Audio [5], and Audio Flamingo 3 [6] provide audio-conditioned text generation. For temporal localization, TimeAudio introduces temporal markers and absolute time-aware encoding [7], while SpotSound interleaves timestamps with audio tokens and uses absent-event queries to suppress hallucinated intervals [8]. SpotSound thus combines event-presence supervision with timestamp prediction. RelTwin targets query-specific selection when inverse relations coexist: both descriptions have acoustic support in the recording, and the requested relation determines which occurrence is the answer.

Compositional audio-language learning. CLAP aligns audio and text in a shared representation space [9]. T-CLAP introduces order-reversed descriptions and a temporal-focused contrastive objective [10]. CompA evaluates event order and attribute binding, and trains CompA-CLAP with composition-aware hard negatives and modular contrastive learning [11]. CoSTALA combines spatio-temporal con-

*Corresponding author: Wei Xu. Email: [to be provided].




Fig. 1. RelTwin binds inverse queries to different windows in one recording. (a) Reused excerpts form two local targets. (b) The dashed enclosure marks training-only computation; shaded diagonal cells identify correct timestamp answers, and up/down arrows indicate training preferences. (c) In a relation-development example, SFT repeats one answer while RelTwin selects the corresponding window for each query. Outlines mark targets; filled bars show predictions, evaluated by set-IoU.

trast with local event-caption alignment and feature consistency [12]. These methods organize audio-text representations around compositional relations, including their local structure. RelTwin places the supervision on complete timestamp answer sequences: their query-conditioned likelihoods determine the preference between competing windows. This trains relational selection in the output space of a generative grounding model.

Hard negatives for temporal grounding. In compositional video grounding, SHINE combines generated hard-negative queries with within-video and cross-query saliency ranking [13]. Its training compares segment relevance as well as query compositions. RelTwin makes complete timestamp answers the candidates for comparison within one recording. Each candidate denotes an entire event sequence with real acoustic support. Under a fixed query, one answer matches the requested order and the other answers the inverse relation; reversing the query exchanges their labels while preserving the audio context. Comparing the length-normalized sequence log-likelihoods explicitly binds each relation query to its corresponding local answer.

3. QUERY-SPECIFIC WINDOW LEARNING

3.1. Two true relations, two different answers

In one recording x , RelTwin arranges the same ESC-50 excerpts [14], A and B , as $A \rightarrow B$ and $B \rightarrow A$. Their local targets W_+ and W_- run from the first excerpt’s onset to the second’s offset, including the gap. Queries q_+ and q_- request the respective orders. Both windows have acoustic support, but W_- is positive for q_- and negative for q_+ ; reversing the query exchanges their labels.

The templates “followed by” and “and then” yield two inverse-query pairs per recording. Earlier-order layouts alternate (Section 5.3). Category rehearsal targets the union of

both occurrences of A or B , retaining category-grounding supervision.

3.2. Contrast timestamp answers during training

We serialize an interval set S in the backbone’s timestamp format: SpotSound-A uses comma-separated spans such as from 1.500s to 11.750s; TimeAudio uses its native temporal tokens. The supervised answer tokens $y(S)$ exclude the prompt and are scored by

$$s_\theta(x, q, S) = \frac{1}{|y(S)|} \sum_t \log P_\theta(y_t | x, q, y_{<t}). \quad (1)$$

The 2×2 matrix M , with entries $M_{ij} = s_\theta(x, q_i, W_j)$ for $i, j \in \{+, -\}$, indexes queries by rows and answers by columns. Its diagonal labels $Y = I_2$ identify the correct query-answer pairs (Fig. 1b). Sequence supervision minimizes $\mathcal{L}_{\text{seq}} = -(M_{++} + M_{--})/2$, rewarding the two correct responses individually. It does not explicitly compare them with the other locally valid window. RelTwin trains the relative preferences $M_{++} > M_{+-}$ and $M_{--} > M_{-+}$ through row-wise probabilities $p_i(j) = \text{softmax}_j(M_{ij}/\tau)$, where $\tau > 0$. The candidate loss is

$$\mathcal{L}_{\text{cand}} = -\frac{1}{2}[\log p_+(+) + \log p_-(-)]. \quad (2)$$

For a separately sampled category example (x_r, q_r, S_r) , the rehearsal loss is $\mathcal{L}_{\text{replay}} = -s_\theta(x_r, q_r, S_r)$. The core objective is

$$\mathcal{L} = \mathcal{L}_{\text{seq}} + \lambda \mathcal{L}_{\text{cand}} + \beta \mathcal{L}_{\text{replay}}. \quad (3)$$

Paired SFT sets $\lambda = 0$; RelTwin uses $\lambda > 0$ with identical relation and rehearsal data. Both directly generate timestamps at inference, without candidates or an additional scorer.

Table 1. Audio temporal grounding (%; higher is better). Baselines: SpotSound Table 3 [8]. Each adapted backbone uses one checkpoint across benchmarks; Δ is the percentage-point gain over its published baseline. $\dagger$: full configuration with exchange consistency and SetPO; SpotSound-A: core candidate objective. Bold marks maxima among listed methods.

Method	SpotSound-Bench			Clotho-Moment			UnAV-100 subset		
	R1@.3	R1@.5	mIoU	R1@.3	R1@.5	mIoU	R1@.3	R1@.5	mIoU
<i>Non-LLM models</i>									
WTATG	53.0	34.5	38.4	12.1	6.3	9.1	53.0	37.0	38.4
AM-DETR	25.5	15.0	19.5	89.8	88.0	80.9	59.0	46.0	42.8
<i>Proprietary models</i>									
Gemini-2.5-Flash	34.8	27.8	25.7	45.7	33.7	36.9	48.0	37.0	35.6
Gemini-2.5-Pro	20.0	17.3	19.7	40.4	32.5	32.5	39.0	35.0	34.6
<i>Open-source models and RelTwin adaptation</i>									
Kimi-Audio	3.0	1.3	3.2	0.7	0.1	0.9	8.0	2.0	5.3
Qwen2-Audio	5.5	1.8	4.4	6.1	1.3	5.7	14.0	4.0	9.7
Audio Flamingo 3	8.8	2.8	7.6	32.9	21.8	22.6	35.0	24.0	25.0
TimeAudio	7.8	1.8	10.2	39.5	24.9	28.6	21.0	7.0	16.0
TimeAudio + RelTwin $\dagger$	24.25	13.25	18.36	45.63	30.92	34.08	44.00	25.00	30.48
Δ vs. TimeAudio	+16.45	+11.45	+8.16	+6.13	+6.02	+5.48	+23.00	+18.00	+14.48
SpotSound-A	76.5	59.5	57.9	93.4	91.0	85.6	86.0	74.0	69.8
SpotSound-A + RelTwin	78.00	62.00	59.43	93.61	91.01	86.68	89.00	79.00	74.03
Δ vs. SpotSound-A	+1.50	+2.50	+1.53	+0.21	+0.01	+1.08	+3.00	+5.00	+4.23

3.3. Joint localization and relation selection

For interval unions, $\text{IoU}_{\text{set}}(S, G) = |S \cap G| / |S \cup G|$ (zero for an empty union). Given predictions $\widehat{W}_i$, define $u_{ij} = \text{IoU}_{\text{set}}(\widehat{W}_i, W_j)$. PairAcc requires $\min(u_{++}, u_{--}) \geq 0.5$. Swap error occurs if $u_{++} \leq u_{+-}$ or $u_{--} \leq u_{-+}$, including ties. JointPairAcc requires PairAcc success without Swap error. We report pair-level percentages; same-answer counts use exact equality of parsed interval lists.

Predicting the union of two disjoint, equal-length targets for both queries passes PairAcc (IoU 0.5) but fails JointPairAcc. Inspired by paired compositional evaluation [15], JointPairAcc distinguishes query-specific selection from coverage.

4. EXPERIMENTAL SETUP

Data and evaluation protocol. Core SpotSound-A adaptation uses 40 ESC-50 classes and 256 recordings: 1,024 relation queries (512 pairs) and 512 rehearsal queries. Ten held-out classes provide 80 recordings, 320 queries, and 160 inverse pairs for relation development, with disjoint classes and original Freesound IDs at adaptation. Evaluation covers all 400 SpotSound-Bench queries over 387 audios, 6,649 Clotho-Moment queries, and UnAV-100’s 100 queries over 77 audios [2]. We report set-mIoU and R1 at IoU 0.3/0.5, clipping intervals to duration. Relation-development and public scores informed method development.

SpotSound-A training. SFT and RelTwin initialize from SpotSound-A on Audio Flamingo 3. LoRA [16] uses rank 8,

$\alpha = 16$, dropout 0.1, and 20.19M parameters. Both run 256 AdamW [17] updates at 5×10^{-6} , zero weight decay, gradient clipping 1, $\tau = 1$, $\beta = 0.5$; RelTwin uses $\lambda = 1$. Each update accumulates one inverse-query pair and one rehearsal query. Initialization, data order, code, environment, and updates are matched. RelTwin computes four gradient-free scores then weighted gradient recomputation; SFT backpropagates through two positive answers. Both include rehearsal.

TimeAudio full configuration. Official TimeAudio, seed 0, receives 384 candidate-training updates with Jensen-Shannon consistency under inverse-query label exchange, then 128 SetPO updates (interval-set listwise preferences with quality-based soft targets). Its 31.76M parameters use rates 2×10^{-6} (LoRA) and 5×10^{-6} (temporal/audio connectors). Event/relation twins mix 3:1; replay mixes Clotho training, synthetic scale/cardinality, TEMPO training, and AudioGrounding training at 4:3:2:1.

Inference. Both backbones use greedy one-beam decoding, capped at 128 new tokens. SpotSound-A uses 16 kHz mono audio and bfloat16; post-development timing tests use frozen matched checkpoints (Section 5.3).

Comparison and statistics. Table 1 compares published baselines; Table 2 compares official SpotSound-A, paired SFT, and RelTwin under shared inference. Official SpotSound-A scores 58.42 mIoU versus published 57.9; RelTwin’s three-seed gain over this re-evaluation is +0.75 (audio-group CI $[-0.24, 1.74]$).

SpotSound-A uses seed 0, selected by SpotSound mIoU from five runs, and matched SFT. Table 2(c) reports seeds

Table 2. SpotSound transfer and synthetic relation-development binding. (b) Matched data and updates; (c) mean $\pm$ sample SD (Section 4). Joint: JointPairAcc; same-answer: inverse-query pairs with identical timestamp answers. Scores are percentages except counts.

Method	SpotSound	Relation development (160 inverse-query pairs)			
	mIoU $\uparrow$	mIoU $\uparrow$	Joint $\uparrow$	Swap error $\downarrow$	Same-answer pairs $\downarrow$
<i>(a) Backbone reference: shared evaluation protocol</i>					
SpotSound-A	58.42	37.64	0.63	65.00	15/160
<i>(b) Matched training comparison</i>					
Paired SFT	58.42	74.58	53.13	40.00	37/160
RelTwin	59.43	88.55	84.38	7.50	1/160
RelTwin – Paired SFT	+1.02	+13.97	+31.25	–32.50	–36 pairs
<i>(c) RelTwin replication: three training seeds</i>					
RelTwin	59.16 $\pm$ 0.26	87.05 $\pm$ 1.43	81.46 $\pm$ 2.53	8.13 $\pm$ 0.63	(1, 0, 0)/160

0/1/2 as mean $\pm$ sample SD. We use 20,000 audio-group bootstrap resamples for 95% CIs; relation tests resample 39 connected source-sharing groups (largest: 18 audios). CIs condition on fixed checkpoints and selection.

5. RESULTS AND ANALYSIS

5.1. Grounding performance across benchmarks

On SpotSound-Bench, Clotho-Moment, and UnAV-100, SpotSound-A + RelTwin reaches 59.43, 86.68, and 74.03 mIoU (Table 1), exceeding its published baseline by 1.53, 1.08, and 4.23 points. Full-configuration TimeAudio reaches 18.36, 34.08, and 30.48, with gains of 8.16, 5.48, and 14.48 points. Each backbone uses one adapted checkpoint across benchmarks.

On the same 400 SpotSound queries, RelTwin gains 1.02 mIoU over matched paired SFT (audio-group CI [0.23, 1.81]): 92 wins, 239 ties, and 69 losses. Its three-seed mean is 59.16 (Table 2c). We next examine responses to inverse queries.

5.2. Timestamp contrasts recover query-specific selection

Paired SFT returns identical answers for 37/160 inverse-query pairs across 22 recordings; RelTwin reduces this to one (Table 2). Of those 37 pairs, 27 become jointly correct: the two predictions recover their respective requested relations. Figure 1 illustrates this query-specific recovery, with each inverse query selecting its corresponding event window.

JointPairAcc increases by 31.25 points (source-group CI [22.35, 40.50]) and Swap error decreases by 32.50. Across all pairs, 51 become jointly correct and one regresses. Relation mIoU increases by 13.97 points (CI [9.84, 17.51]), connecting selection with overlap.

Across three RelTwin seeds, JointPairAcc averages 81.46 and identical-answer counts are 1, 0, and 0, reproducing the reduction in same-answer collapse.

5.3. Binding gains persist after timing intervention

To test layout-specific silence, we recompose 80 development audios with 1.5 s lead-in, 0.25 s within-window gaps, 3 s between windows, and 2 s tail, preserving waveforms, order, and queries. The frozen matched models evaluate one template (160 queries). Timing changes for 40 BA-first recordings; 40 AB-first recordings remain waveform-identical.

On changed recordings, RelTwin reaches 87.50 JointPairAcc against SFT’s 67.50 (+20.00; CI [2.94, 33.33]); originally they scored 85.00 and 42.50. RelTwin retains a 20-point selection advantage under the altered timing, supporting query–window binding across both silence layouts. Across all 160 queries, SFT and RelTwin score 81.34 and 90.65 mIoU.

5.4. Why locally valid answers matter

Shared excerpts hold event identity fixed while inverse queries exchange the correct window. Detecting both categories supports both candidates; selecting the requested order requires query-conditioned comparison. RelTwin’s reduced same-answer collapse and improved JointPairAcc capture this distinction in the controlled binding test. Timestamp-answer supervision preserves direct generation at inference.

6. CONCLUSION

RelTwin fine-tunes grounding models through locally valid timestamp contrasts: acoustically supported windows compete according to the query. The adapted SpotSound-A model exceeds its published mIoU on three benchmarks. Matched diagnosis and timing intervention demonstrate improved query–window binding. These results support explicit supervision of which occurrence answers the query while preserving direct timestamp generation at inference.

7. ACKNOWLEDGMENTS

Wei Xu personally funded GPU computing. No institutional or corporate funding was received. The authors declare no relevant financial or nonfinancial conflicts of interest.

Compliance with Ethical Standards. We reuse public ESC-50 [14], SpotSound [8], Clotho-Moment, UnAV [2], TEMPO, and AudioGrounding data, including programmatically constructed sequences. No new recordings, participant recruitment, human or animal experiments, or listening studies were undertaken.

8. REFERENCES

- [1] Xuenan Xu, Heinrich Dinkel, Mengyue Wu, and Kai Yu, “Text-to-audio grounding: Building correspondence between captions and sound events,” *arXiv preprint arXiv:2102.11474*, 2021.
- [2] Hokuto Munakata, Taichi Nishimura, Shota Nakada, and Tatsuya Komatsu, “Language-based audio moment retrieval,” *arXiv preprint arXiv:2409.15672v3*, 2025.
- [3] Xuenan Xu, Ziyang Ma, Mengyue Wu, and Kai Yu, “Towards weakly supervised text-to-audio grounding,” *arXiv preprint arXiv:2401.02584*, 2024.
- [4] Soham Deshmukh, Benjamin Elizalde, Rita Singh, and Huaming Wang, “Pengi: An audio language model for audio tasks,” in *Advances in Neural Information Processing Systems*, 2023, vol. 36.
- [5] Yunfei Chu, Jin Xu, Qian Yang, Haojie Wei, Xipin Wei, Zhifang Guo, Yichong Leng, Yuanjun Lv, Jinzheng He, Junyang Lin, Chang Zhou, and Jingren Zhou, “Qwen2-Audio technical report,” *arXiv preprint arXiv:2407.10759*, 2024.
- [6] Arushi Goel, Sreyan Ghosh, Jaehyeon Kim, Sonal Kumar, Zhifeng Kong, Sang gil Lee, Chao-Han Huck Yang, Ramani Duraiswami, Dinesh Manocha, Rafael Valle, and Bryan Catanzaro, “Audio Flamingo 3: Advancing audio intelligence with fully open large audio language models,” *arXiv preprint arXiv:2507.08128*, 2025.
- [7] Hualei Wang, Yiming Li, Shuo Ma, Hong Liu, and Xiangdong Wang, “Listening between the frames: Bridging temporal gaps in large audio-language models,” in *Proceedings of the AAAI Conference on Artificial Intelligence*, 2026, vol. 40, pp. 26233–26241.
- [8] Luoyi Sun, Xiao Zhou, Zeqian Li, Ya Zhang, Yanfeng Wang, and Weidi Xie, “SpotSound: Enhancing large audio-language models with fine-grained temporal grounding,” *arXiv preprint arXiv:2604.13023v2*, 2026.
- [9] Yusong Wu, Ke Chen, Tianyu Zhang, Yuchen Hui, Taylor Berg-Kirkpatrick, and Shlomo Dubnov, “Large-scale contrastive language-audio pretraining with feature fusion and keyword-to-caption augmentation,” in *IEEE International Conference on Acoustics, Speech and Signal Processing*, 2023.
- [10] Yi Yuan, Zhuo Chen, Xubo Liu, Haohe Liu, Xuenan Xu, Dongya Jia, Yuanzhe Chen, Mark D. Plumbley, and Wenwu Wang, “T-CLAP: Temporal-enhanced contrastive language-audio pretraining,” *arXiv preprint arXiv:2404.17806*, 2024.
- [11] Sreyan Ghosh, Ashish Seth, Sonal Kumar, Utkarsh Tyagi, Chandra Kiran Reddy Evuru, Ramaneswaran S, S Sakshi, Oriol Nieto, Ramani Duraiswami, and Dinesh Manocha, “CompA: Addressing the gap in compositional reasoning in audio-language models,” in *International Conference on Learning Representations*, 2024.
- [12] Peiwei Ren, Jinbo Hu, Fang Kang, Shan Liang, and Yin Cao, “CoSTALA: Compositional spatio-temporal audio-language alignment via multi-grain hierarchical contrastive learning,” *arXiv preprint arXiv:2608.24374*, 2026.
- [13] Zixu Cheng, Yujiang Pu, Shaogang Gong, Parisa Kordjamshidi, and Yu Kong, “SHINE: Saliency-aware hierarchical negative ranking for compositional temporal grounding,” in *European Conference on Computer Vision*, 2024.
- [14] Karol J. Piczak, “ESC: Dataset for environmental sound classification,” in *Proceedings of the 23rd ACM International Conference on Multimedia*, 2015, pp. 1015–1018.
- [15] Tristan Thrush, Ryan Jiang, Max Bartolo, Amanpreet Singh, Adina Williams, Douwe Kiela, and Candace Ross, “Winoground: Probing vision and language models for visio-linguistic compositionality,” in *IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 2022.
- [16] Edward J. Hu, Yelong Shen, Phillip Wallis, Zeyuan Allen-Zhu, Yuanzhi Li, Shean Wang, Lu Wang, and Weizhu Chen, “LoRA: Low-rank adaptation of large language models,” in *International Conference on Learning Representations*, 2022.
- [17] Ilya Loshchilov and Frank Hutter, “Decoupled weight decay regularization,” in *International Conference on Learning Representations*, 2019.