

(a) Shared excerpts



Inverse queries exchange the correct answer

q_+ : dog followed by rooster $\rightarrow W_+$

q_- : rooster followed by dog $\rightarrow W_-$

(b) RelTwin: two-stage fine-tuning

1 Query-conditioned competition

Score complete timestamp answers

	$y(W_+)$	$y(W_-)$
q_+	$s_{++} \uparrow$	$s_{+-} \downarrow$
q_-	$s_{-+} \downarrow$	$s_{--} \uparrow$

Candidate CE + label-exchange JS

2 Set-level preferences (SetPO)

Compare coverage, count, and boundaries



Quality scores $\rightarrow$ soft preference targets

TRAINING

Inference

Audio + query $\rightarrow$ Adapted backbone $\rightarrow$ Timestamps

(c) Query-specific selection

Inverse-query grounding example

Paired SFT

+ Candidate

